

N. OOZEER, M. BIETENHOLZ & S. GOEDHART

INTRODUCTION TO CASA: A KAT-7 DATA REDUCTION GUIDE

... YOU CANNOT ACQUIRE EXPERIENCE BY MAKING EXPERIMENTS. YOU CANNOT CREATE EXPERIENCE. YOU MUST UNDERGO IT.

ALBERT CAMUS.

Contents

1	<i>Why do you want to learn CASA?</i>	5
1.1	<i>How to use this document</i>	6
1.2	<i>Conventions used</i>	7
1.3	<i>Methodology</i>	8
2	<i>Let us get serious</i>	9
2.1	<i>Overview of Calibration</i>	9
2.2	<i>STEP 1... Converting the hdf5 file to a MS</i>	9
2.3	<i>STEP 2... Loading and inspecting the data</i>	10
2.4	<i>STEP 3... Flagging</i>	16
2.5	<i>STEP 4... Calibration</i>	22
2.6	<i>STEP 5... Splitting the data</i>	36
2.7	<i>STEP 6... CLEAN aka Deconvolution</i>	38
3	<i>Spectral line calibration and imaging</i>	43
3.1	<i>Inspection and flagging</i>	43
3.2	<i>Set up models for the calibrators</i>	45
3.3	<i>Calibration</i>	46
3.4	<i>Velocity rest frames</i>	53
3.5	<i>Continuum subtraction</i>	57
3.6	<i>Finally, on to imaging</i>	59

4 *Troubleshooting and Support* 63

4.1 *Getting help* 63

Bibliography 65

Why do you want to learn CASA?

For quite some decades, radio astronomy data reduction has been performed using AIPS, Miriad and other packages. The coming online of new and substantially enhanced radio interferometers like ALMA, LOFAR and the JVLA has driven the development of a new software package for interferometric data named CASA. CASA is better suited to the complexity and data-volume from these instruments than previous existing packages. CASA was developed by NRAO, and was built on the Aips++ codebase. It seems to be becoming the software suite of choice for radio interferometric data. KAT-7 is a small interferometer array and this tutorial presents a guide to reducing the Karoo Array Telescope (KAT-7) data using CASA.

Producing an image from interferometer data is not straightforward: an interferometer does not form an image in the way an optical telescope does, but rather the observations are made in a spatial frequency space, called the u - v or visibility plane, which is essentially a Fourier transform of the image plane. So, in order to make an image, the interferometer measurements must be Fourier transformed.

Due to the nature of the interferometer, the (u, v) plane is not fully sampled. The presence of unsampled regions in the (u, v) causes artifacts to appear in the image plane. The actual measurements made by the interferometer are called visibility measurements. In particular, the geometry of the telescopes in relation to the direction to the source determines the instantaneous u - v coverage of an interferometer. Often the source is observed for some time, and as the earth and the telescopes on it rotate with respect to the source, the pattern of (u, v) coverage also rotates, allowing more parts of the (u, v) plane to be sampled. The Fourier transform of this sampling function is the point-spread function of the instrument, called the “dirty beam” in radio interferometry.

An interferometer is inherently a somewhat complex instrument, and modern radio interferometers all produce multi-layered data sets which have visibility measurements from all the different baselines



Figure 1.1: One of the KAT-7 dishes.

forming the interferometer, possibly of several different polarization products, and in addition, they often have a large number of independent frequency channels. Depending in the science aims, these independent channels can be combined to form a single-channel image, or they can be used to produce a multi-channel image cube.

However, before the visibility measurements can be sensibly Fourier transformed to the image plane, they must be calibrated. Each telescope has a complex gain which is not known a priori, but must be determined from the observations. Often specific sources with known properties, called calibrator sources, are observed with the express purpose of determining these gain factors. The visibility measurements on any baseline are affected by the complex gains of both telescopes making up that baseline. The first stage of data processing in aperture synthesis is to determine the amplitude gain and the phase-shift due to each individual telescope (and the atmosphere above it), which together form the complex gain for that telescope. Although an interferometer is designed to keep these gains as stable as possible, they do vary with time, in particular the phase part.

Data sets often also contain some bad or corrupted data, most commonly due to radio frequency interference (RFI¹) or to malfunction of some part of the antenna electronics. This bad data should be removed, and the process of removing them is known as flagging. If the calibration is to be determined from observations of calibrator sources, then any bad data for the calibrator sources should ideally be removed before determining the calibration solutions. For other sources it is generally easier to identify bad data after the calibration is applied.

¹ http://pos.sissa.it/archive/conferences/107/001/RFI2010_001.pdf

Once the visibility data are properly calibrated and edited, they can then be Fourier transformed to make an image. However, due to the incomplete sampling in the image plane, the image so produced is the convolution of the sky brightness with the point spread function of the interferometer, which is called the “dirty beam”. The second part of the data reduction is to try and correct the image, to the extent possible, for the effects of this incomplete sampling. Unlike data at some other wavelengths, radio data can require hours to days of reduction to achieve the highest dynamic range. Ideally, the image meets the constraints of physical consistency as well as plausibility. This document will attempt to explain the various stages for getting radio images from KAT-7 data.

1.1 *How to use this document*

This document assumes that you are conversant with Linux and Python. If not, then the book “SAMS teach yourself Linux in 10 minutes” will be a good start even though you will take more than 10 min-

utes to read the whole book. Reading “Python for Dummies” will also help. We will also not cover all the fundamentals of radio astronomy, but from time to time we shall try to discuss some concepts before proceeding so that the user is aware of what we are doing. For a thorough introduction to the fundamentals of radio astronomy, we recommend the NRAO Essential Radio Astronomy online course (<http://www.cv.nrao.edu/course/astr534/ERA.shtml>).

The procedures consist of various steps and each step generally depends on the previous ones. So avoid skipping to the next step if you have warnings and errors, without clearing them. We will note those steps that are independent of the preceding ones.

1.2 Conventions used

We have tried to use the following conventions through this document:

- A typewriter font is used for CASA tasknames and variables in the main text, and for CASA input and outputs given in a red box, for example:

CASA command:

```
prefix = 'CirX1'
msfile = prefix+'.ms'
```

- The Return key is synonymous with the ENTER key or ↵



Tips: lead to shortcuts that we find useful



Cautions: Avoid pitfalls



Comments: explain the relevant procedure



Pause: Time for a break

1.3 Methodology

In this document we will review the data reduction procedures that we will use during the in-house workshop on CASA.

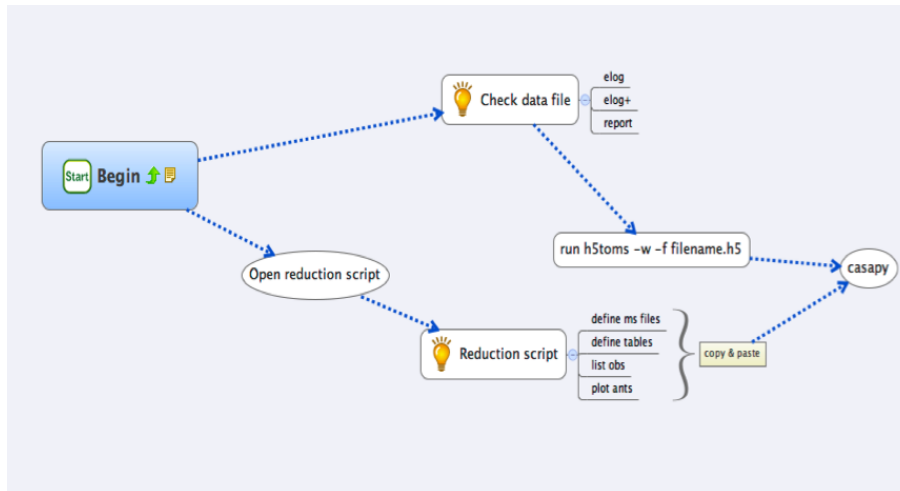


Figure 1.2: This mindmap shows the routes for getting your data into CASA

Before you begin in this long journey of data reduction, make sure you have enough time to spend. It can get very crazy; get lot of snacks, coffee and put THE sign (Figure 1.3) in a place clearly visible to the rest of the world. Basically we expect for a newbie to spend around 3-5 days to successfully reduce a KAT-7 dataset. However, this documentation will allow you to get a clean image in a day, the target is within 10 hours of CONTINUOUS work.

It is not that hard but requires lot of attention and concentration. The process will become quicker after you've worked with a few data sets and have become familiar with it.

OK let us get serious now.



Python is case sensitive (as is Linux). Always ensure you using the correct case.



Figure 1.3: THE sign

2

Let us get serious

2.1 Overview of Calibration

The visibility measurement made by the interferometer are complex values, having either real and imaginary parts, or an amplitude and a phase. The calibration process consists mostly of finding a set of time and frequency dependent complex gain factors, by which our visibilities need to be multiplied.

We develop our complex gain solution in stages. Some parts vary with time but not so much with frequency, for example transmission through the atmosphere causes both phase shifts and attenuation. However, the variation of the atmospheric part of the complex gain across the observing band is usually small. Therefore, a combined solution for all the different frequency channels of observations is usually possible. The bandpass response of the instrument, by contrast, varies across the frequency band, but is usually quite stable in time, so that it needs to be solved for only once in an observing run.

So our whole set of complex gain factors is usually broken up into several pieces, each of which is solved for separately. This division is not necessarily unique, in the end all that matters is the complex product of the individual parts of the calibration.

2.2 STEP 1... Converting the hdf5 file to a MS

Note that for the data set we are using for the tutorial, this step has already been done, so you can skip ahead to the next section. If you need to do this step, go to the KAT-7 archive to search for your file. Then check the eLog to see if the telescope operator entered any information during the observations, and take note of any if they did. Then, proceed to converting the archive file to a Measurement Set (ms). An ms is the form in which CASA will work on your data. Also remem-

ber at the end of the data reduction to add an entry in the eLog of the results. Convert the hdf5 archive file to an ms using the in-house script `h5toms` as follows:

linux command:

```
h5toms.py -f --no-auto -C '199,799' --flags='all' -r ant5 myfile.h5
```

The `-f` flag causes `h5toms.py` to produce a full-polarization output file, while the `no-auto` causes `h5toms` to load only cross-correlations, but not the auto-correlations. The `-C 199,799` flag causes `h5toms` to load only channels 199 to 799 (inclusive) of the original 1024 channels. The ~ 200 channels at either edge of the band are known to be bad, so we do not load them. We have also applied the online flagging using the `--flags='all'`. The `-r ant5` causes it to use antenna 5 as the reference antenna (although you can change this later if required).

The output will be a directory `myfile.full_pol.ms`. This directory is your CASA measurement set. More options are available for `h5toms.py`¹. Once the data has been successfully converted into an ms, one can start CASA from the computer `kat-imager2` (assuming you are running on the SKA Pinelands server), and then proceed with the data reduction.

¹ `h5toms.py -h` lists the additional options e.g. for `h5toms`

It's probably a good idea to rename your ms to something a little more descriptive at this stage.

linux command:

```
mv myfile.full_pol.ms School_data_av.ms
```

Of course if you're looking at real data its probably a good idea to use a name for the ms which reflects what the observations are.

2.3 STEP 2... *Loading and inspecting the data*

It is probably best to create a specific directory which will contain the working CASA data. We will use the directory "`my_working_directory`".

linux command:

```
mkdir my_working_directory
```

The first main step is to move the ms into your chosen working directory, and then also to make the working directory your current directory. For the workshop, you may have downloaded the file `School_data_av.ms.tar.gz`, which you would unpack with the linux command:

linux command:

```
tar -zxf School_data_av.ms.tar.gz
```

It should produce a directory called `School_data_av.ms`. This directory is your CASA measurement set.

#

Locate the ms directory from the previous step and move it into the working directory.

```
mv myfile.full_pol.ms my_working_directory
cd my_working_directory
```

#

Although CASA measurements sets (ms) are ordinary linux directories, and can be renamed, moved, deleted or copied like any other directories, there are probably restrictions you should observe:

- Do not move, delete etc. directories from the operating system *while casapy is running*, as casapy does some ms caching, and will likely get confused. You can remove an ms from within casapy by typeing e.g., `rmtables('crappy-data.ms')` and it will be fine, or you can just exit casapy and then delete it.
- Do not alter the *contents* of the ms. While renaming or moving the ms directory is fine, its probably not a good idea to add or delete any files to the ms, or to alter or rename any of the files contained within it.

Most of the data reduction will be done from within casapy, which provides an interactive shell (essentially the equivalent of interactive python or ipython). From the Linux prompt you will therefore first run casapy. You will then give the commands to run various CASA programs from within casapy. In casapy you can also create variables, assign values to them, and then subsequently use them at any time during the processing. It is generally convenient to define some variables to hold the names of ms's and tables. The sequence of casapy commands can then be re-used for a different data set by merely changing the values assigned to these variables.



Do NOT abort running programs by pressing Ctrl+C as CASA has a bug and will screw up all your visibilities. Rather wait for the command to be over and then undo the wrong steps.

In fact, CASA is designed to be highly script-able, and sequences of commands could be placed in a script and run automatically. In this session, however, we shall be using step-by-step commands.



Never delete the ms file unless you have finished splitting your relevant data and are sure you do not need the ms anymore.

Below we define a variable to hold the ms (directory) name but without the “ms” extension. We will use this variable to refer to the ms and to various associated tables. The ms we are using in the tutorial is called Cirx1_school.ms, so we define the following variables in casapy:

CASA command:

```
prefix = 'Cirx1_school'  
msfile = prefix+'.ms'
```

We now define variables that hold the names of various tables that will be used during the reduction. The table names could be specified directly, but if you do it by means of variables, the subsequent commands we use can be copied and pasted and re-used for some other reduction with the only modification required being the re-definition of these variables.

CASA command:

```
bandpass_table0 = prefix + '_spw0.B0'  
gain_table0 = prefix + '_spw0.G0'  
gain_table1 = prefix + '_spw0.G1'  
flux_table1 = prefix + '_spw0.fluxscale1'
```

We chose antenna ant5 as our reference antenna when we first converted the data-set to an ms, because it seems to have more shielding for RFI from the pedestal. Since it is best to use a consistent reference antenna throughout, we again define a variable whose value is the current reference antenna.

CASA command:

```
reference_antenna= 'ant5'
```

Next, we use the task `listobs` to print out a summary of the observations in the ms (similar to the AIPS task `LISTR` with `OPTY='SCAN'`):

CASA command:

```
listobs(vis=msfile)
```

CASA output:

(we have removed the leading part of the lines of the listobs output for clarity.)

```
=====
MeasurementSet Name: /home/michael/casa-workshop2012/CirX1.ms      MS Version 2
=====

Observer: lindsay      Project: 20120701-0006
Observation: KAT-7
Data records: 71904      Total integration time = 56296.5 seconds
Observed from 01-Jul-2012/13:57:43.4 to 02-Jul-2012/05:35:59.9 (UTC)

ObservationID = 0      ArrayID = 0
Date      Timerange (UTC)      Scan  FldId FieldName      nRows  Int(s)  SpwIds
01-Jul-2012/13:57:43.4 - 13:59:27.9      1      0 PKS 1934-638      168    14.1    [0]
      14:00:13.4 - 14:04:58.4      2      1 Circinus X-1      420    14.8    [0]
      14:05:28.4 - 14:06:21.4      3      2 PKS 1613-586      105    11.6    [0]
      14:06:44.5 - 14:11:30.5      4      1 Circinus X-1      420    14.8    [0]
      14:12:00.5 - 14:12:45.5      5      2 PKS 1613-586      84     14.5    [0]

... (we omit the bulk of the detailed scan listing here)

      05:29:53.8 - 05:30:38.9      328      2 PKS 1613-586      84     15      [0]
      05:30:56.9 - 05:31:41.4      329      2 PKS 1613-586      84     14.5    [0]
      05:31:58.9 - 05:32:43.9      330      2 PKS 1613-586      84     15      [0]
      05:33:01.9 - 05:33:46.9      331      2 PKS 1613-586      84     14.8    [0]
      05:34:04.9 - 05:34:49.4      332      2 PKS 1613-586      84     14.8    [0]
      05:35:06.9 - 05:35:59.9      333      2 PKS 1613-586      105    12      [0]
(nRows = Total number of rows per scan)

Fields: 3
ID  Code Name      RA      Decl      Epoch  nRows
0   T   PKS 1934-638  19:39:25.03058 -63.42.45.6999 J2000 2709
1   T   Circinus X-1  15:20:40.85208 -57.10.00.1048 J2000 52542
2   T   PKS 1613-586  16:17:17.89966 -58.48.07.8890 J2000 16653

Spectral Windows: (1 unique spectral windows and 1 unique polarization setups)
SpwID #Chans Frame Ch1(MHz)  ChanWid(kHz) TotBW(kHz) Corrs
0      19 REST  1938.21094  -12500      234765.625 XX XY YX YY

The SOURCE table is empty: see the FIELD table
Antennas: 7:
ID  Name Station  Diam.  Long.      Lat.
0   ant1 ant1     12.0 m +021.24.39.4 -30.33.10.2
1   ant2 ant2     12.0 m +021.24.41.9 -30.33.09.1
2   ant3 ant3     12.0 m +021.24.38.6 -30.33.09.1
3   ant4 ant4     12.0 m +021.24.37.7 -30.33.09.1
4   ant5 ant5     12.0 m +021.24.37.1 -30.33.10.0
5   ant6 ant6     12.0 m +021.24.36.2 -30.33.12.5
6   ant7 ant7     12.0 m +021.24.35.2 -30.33.07.5
```

It is a good idea to save the listobs output for future reference, as it gives the basic structure of the observations, indicating for example which source was observed at what time. You can either copy the listobs output from the casapy log file, or rerun listobs with extra parameter listfile=[insert the name of the textfile you want to save the output here].

What can you deduce from this listing of the data set?

- The observations started at 14:00 on 01-Jul-2012 and continued till 04:15 on 02 Jul 2012 (UTC times)
- Three sources were observed, with field names “PKS 1934-638”, “Circinus X-1” and “PKS 1613-586” and source IDs 0, 1, and 2 respectively.
- The reference frequency 1938.21094 MHz; One spectral window, SPW, was used, with SpwID = 0. This SPW had 19 frequency channels (note that this data-set has already been averaged in frequency, raw data KAT-7 data sets typically have 1024 channels)
- The number of antennas taking part was 7

The science target in these observations was Circinus X-1 (ID=1), with PKS 1934-638 and PKS 1613-586 (ID 0 and 2 respectively) being the calibrators. The source PKS 1934-638 will be used as a flux density calibrator, while PKS 1613-586 will be used to calibrate the antenna phases.

In CASA, it is the `field` parameter which is generally used to specify the source, and one can use either the source name (e.g., `field = 'Circinus X-1'`) or the Source ID (e.g., `field = '2'`) to specify a particular source. Note that the value of `field` is a string, and therefore enclosed in quotes (‘ ’) even when it is specifying the numeric source ID.

CASA can automatically set the flux densities for the flux-density calibrator sources, but they need to have field names in the correct format for this to work. In this case, our flux-density calibrator source has the form of the name recognized by CASA “PKS 1939-638”. However, if your flux-density calibrator source does not have a name of the recognized form, for example, “3C 48”, then you need to use `browsetable` to replace it with, e.g., “3C48”. In a future version of the `h5toms` script this will be fixed automatically. `browsetable` will produce a GUI where you can edit the `ms` tables. Check the link <http://casa.nrao.edu/docs/UserMan/UserMansu193.html> for details.

It is good practice to plot the antenna positions and check that everything being read properly and your antennas are where they should be. This is done using:

CASA command:

```
plotants(vis=msfile)
```

From the above antenna plots, we can see that ant4 was not used in this observation and also that the plots show the expected positions of the KAT-7 antennas.

2.4 STEP 3... *Flagging*

The majority of interferometric data sets contain some corrupted data, with the most common reason being failures of an antenna or some of its associated electronics and radio frequency interference (RFI). The process of editing out this data is generally called “flagging”, and is usually the most boring and tedious part of the data reduction process. Flagging is somewhat of an art and there are no hard rules. However, most modern instruments, KAT-7 in particular, produce fairly clean data (at least at frequencies above 1 GHz), so once you have removed the data from any malfunctioning antennas, you should not have to flag any more than 10% of your data, and often much less than that. If possible you should flag either all visibilities for some time interval, or all by antenna, rather than flagging individual baselines. The reason for this is twofold: firstly, its more efficient, and secondly its less biased. Although different interactive flagging tools are available, for the initial flagging the best approach is usually to first make several passes through the data, using various plotting and listing tools, to identify the bad data. Once you have identified some sections of bad data, you then flag them explicitly using `flagdata`. This allows you to keep a detailed record of what was flagged, and to easily redo the flagging if required. Going through your visibility data set point by point and picking out individual points is rarely productive.

The following tasks are available for flagging data in CASA:

- `flagmanager`: save and manage versions of the flagging entries in the ms
- `flagautocorr`: noninteractive flagging of auto-correlations
- `plotms`: interactive XY plotting and flagging of visibility data
- `plotxy`: interactive XY plotting and flagging of visibility data - note that `plotxy` is slower than `plotms` and will eventually be phased out
- `viewer`: can display (as a raster image) ms data, with some editing capabilities
- `flagdata`: This is the new non-interactive flagging (and unflagging) of specified data in CASA 3.4 onwards. It has other flagging options such as antenna shadowing, elevation flagging, etc.

We will mostly use the task `plotms` to examine the data. It is a very flexible tool that allows you to interactively plot your visibility data in a wide variety of ways. We will use it to identify the bad data. We will then use the task `flagdata` to actually flag it. Although data *can* be flagged directly in `plotms`, we suggest that it is better to actually identify which antennas, baselines and times are responsible for bad data before flagging them until you have some experience. Flagging data *only* on the basis of e.g., low visibility amplitude can lead to serious biases: if you expect a flux density of 1 Jy and you flag all the visibilities outside the range 0.9 to 1.1 Jy, you would likely find the remaining visibilities suggested a source of 1 Jy. However, you're not learning about the sky, you're learning about the flagging you already did!

Calibrator sources are much easier to flag than program sources, because we largely know what kind of source they are, and therefore what to expect from the visibility measurements. In particular, calibrators are generally point-like, which means the true visibility amplitude is independent of baseline (and of time), and the phase should be 0° . Unfortunately, all we have available at this point is the *uncalibrated* visibilities. However, for an instrument like KAT-7, where all the dishes, and thus all the baselines have approximately equal sensitivity, even for the uncalibrated data for a point-like source, we can say this: the (uncalibrated) visibility amplitude on all baselines should be approximately the same, say within a factor of 2 or so, and the (uncalibrated) visibility phases should be fairly stable, although they do not need to be near 0° .

The two most common failures to identify are probably “dead” or malfunctioning antennas, and channels which are corrupted by RFI, which is usually very narrow bandwidth. Even though automatic flagging has already been done by `h5toms.py`, and may well have removed most of the strong RFI, you should nonetheless always examine your data.

However, first let us remove some bits data which is known to be bad. The KAT-7 correlator is known to produce some visibilities with spurious amplitudes of exactly 0. We will use `flagdata` to delete these. We will also use a second invocation of the `flagdata` to flag any data taken at antenna elevations below 10° . Data taken at very low antenna elevations is often unreliable and noisy, so unless much of your data is taken at low elevations, its often easier to just delete the small fraction with antenna elevations below about 10° .

CASA command:

```
flagdata(vis=msfile,mode='clip',field='',
         clipzeros=True, flagbackup = False)
```

```
flagdata(vis=msfile,mode='elevation',
         lowerlimit=10.0, flagbackup=True)
```

Note that in “elevation” mode, the elevation limits provided are those of the *good* data, rather than those of the data to be flagged, so in this case any data below an elevation of 10° will be flagged. We do not need to set an upper elevation limit.

Now let's plot the remaining visibilities. Phase stability, both in frequency and time, is probably the most important thing to check first: unless the visibility phases for the calibrator sources are reasonably stable in time and across the different channels, calibration will likely prove impossible.

Each visibility point is associated with one baseline, that is some *pair* of antennas. However, instrumental failures are almost always antenna-dependent, that is, due to a failure of some particular antenna, and therefore *all* the baselines involving the bad antenna will have bad visibility data. When identifying data to flag, it is best to try first to establish whether it is some antenna that is bad, and if so flag all the data to that antenna. Some bad data, however, in particular that due to RFI, may not be antenna-dependent, and may have to be flagged individually per baseline.

For starters, let us look at the phase as a function of frequency on a per-baseline basis. We will start with just one example scan of our strongest calibrator, PKS 1934-638, pick one scan from the listing above, somewhere near the middle of the run. In this case we will pick scan ID = 57, occurring between times 17:01:16.3 and 17:03:01.3.

Before you carry on, do spend some time getting to know the GUI behind `plotms` and how to use it:

http://casaguides.nrao.edu/index.php?title=Data_flagging_with_plotms.

For axis definition in `plotms` check this link :

http://casaguides.nrao.edu/index.php?title=What%27s_the_difference_between_Antenna1_and_Antenna2%3F_Axis_definitions_in_plotms.

Now actually run `plotms` to bring up the plot window and get a first look at part of our visibility data. Note that the `iteraxis='baseline'` tells `plotms` to plot each baseline individually. Pressing the green arrow at the bottom of the `plotms` interface will allow you to step through the individual baseline plots interactively.

CASA command:

```
plotms(vis=msfile, xaxis='channel', yaxis='phase',
       correlation='XX,YY', scan='57', field='PKS 1934-638',
       iteraxis='baseline', coloraxis = 'corr',
       plotrange=[0,0,-180,180])
```

Your screen will look similar to this:

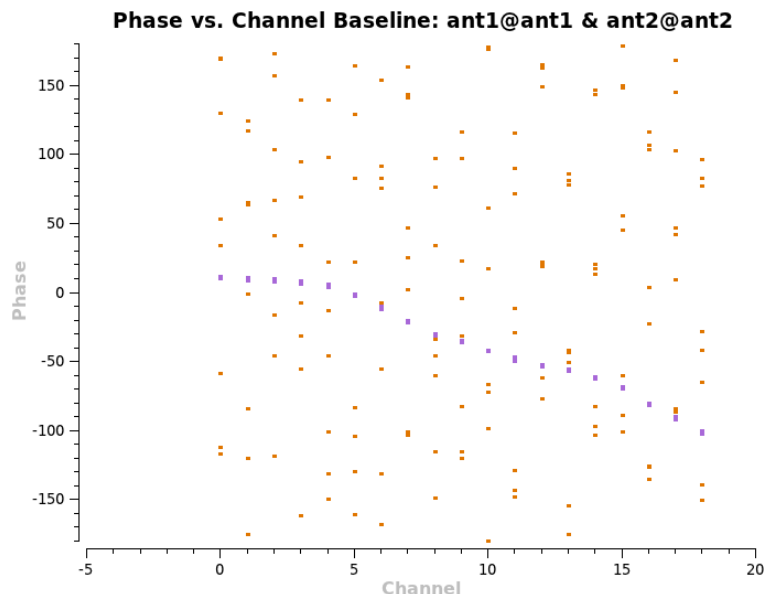


Figure 2.1: The output of plotms showing phase versus channel for one baseline.

The two different colours represent the two correlations or polarization products (XX, YY) which we have plotted. Note that the purple points, which correspond to XX, hang together well, sloping slightly from channel 0 to channel 19. They are fine. The tan-coloured points for the other polarization, however, are scattered randomly, and will not be usable. This plot is for the baseline between Ant1 and Ant2. Hitting the right arrow button at the bottom will show the plot for the next baseline Ant1, Ant3, where both polarizations look fine. The next plot, Ant1, Ant4 again shows the XX points are fine, while the YY points show strong variation across the band.

These effects are almost certainly due to an *antenna*, rather than a single baseline, and the trick is to identify which of the antennas is the one that is affected. By stepping through the baseline you will see that the YY phases for all baselines involving ant2 look quite random, while those involving ant4 are somewhat less random, but nonetheless show large variation from channel to channel. In other words, at least for this scan, there is a problem with the phase stability for the YY

polarization of Ant2 and Ant4.

Let us check for phase stability against time. We will only plot one channel by specifying `spw='*:9'` for this purpose. (Seeing as the phase will vary both with time and channel, its hard to determine the phase stability by looking at all channels and times simultaneously, that's why we've taken the approach of first taking a "slice" in time, and looking at only one scan but all channels. Now we take a "slice" in frequency, and look at only one channel but at all times).

CASA command:

```
plotms(vis=msfile, xaxis='time', yaxis='phase',
       correlation='XX,YY', spw='*:9', field='PKS 1934-638',
       iteraxis='baseline', coloraxis = 'corr',
       plotrange=[0,0,-180,180])
```

Again you will see a similar pattern. The visibility phases for XX are quite stable in time, showing only small variations from scan to scan and within each scan (mostly, the ~ 10 or so visibility points within a scan are so close together that the whole scan appears as a single point in the plot). However, again, the visibility phases for the YY polarization (tan colour) for ant2 and ant4 are quite unstable.

Since we have shown that the visibility phases for YY, ant2 and ant4 are quite unstable both in time and across channels, we will delete these points. Note, however, that so far, we have only examined the data for one source, namely PKS 1934-638. Perhaps the data for the other sources are fine! Before throwing out all the data, its worth checking. In the `plotms` window, change "field" to PKS 1613-586 and hit "plot". There are more data for this source, but you will see the exact same pattern. Since the data for YY, ant2 and ant4 are bad for both our calibrator sources, we have no choice but to flag all these data, including that for our target source Circinus X-1. Very likely the target source data would also be bad, but in any case it cannot be calibrated so we cannot use it. You *could* plot also the visibility data for Circinus X-1, however since its not known to be a compact source we do not know what to expect from the visibility phases. They may vary strongly because the source is complex. It is only for calibrator sources, where the actually visibility phase should be constant that we can judge stability by plotting the visibility phases as we have done.

At present CASA has the limitation that it cannot calibrate one of the XX, YY polarizations when the second one is missing, so we will actually have to flag all the polarizations for ant2 and ant4, even though the XX polarization data does not have any problem. (There are four

total correlations, XX, YY, XY, YX. We will deal only with first two, whose sum represents the unpolarized intensity. The calibration of the polarized parts, XY, YX involves extra steps, and is not yet implemented in CASA for linearly polarized feeds like those of KAT-7). Having identified the bad data, we once again turn to `flagdata` to actually flag it. We do not specify either field or correlation so `flagdata` flags the requested data for all fields and correlations.

CASA command:

```
flagdata(vis=msfile, antenna='ant2,ant4',
        flagbackup=True)
```

#

Referring to an antenna can be kind of confusing in CASA. Like sources, CASA can refer to antennas by their IDs or numbers, or by their names. The confusing thing is that the antenna names often look very similar to IDs, with KAT antennas having names like, e.g., “ant1”. However, the antenna IDs start with 0, while the names start with the number 1. You have to keep track of whether you are referring to the antenna by ID or by name. In this case, the antenna ID=0 has the name “ant1” (as can be seen from the `listobs` output). So it is antennas ant2 and ant4, also known as ID = 1 and 3, which have the bad data in this case.

If you now repeat the above `plotms` commands, you will see that the visibility phases are now nicely coherent for all the plotted baselines, and the baselines we have flagged no longer show up.

Now let's turn to the visibility amplitudes. Again we make use of the fact that for continuum calibrator sources, the expected values are known and should be the same on all baselines, and vary only slightly across the channels. This time we can plot all the baselines together, rather than iterating through them as we did before:

CASA command:

```
plotms(vis=msfile, xaxis='channel', yaxis='amp',
       correlation='XX,YY', field='PKS 1934-638',
       coloraxis='corr')
```

You will see that for most channels there is only a relatively small variation of the visibility amplitude with channel for PKS 1934-638 (note the amplitude scale on the left). This small variation will be corrected later with `bandpass`. However, channel ID=6 shows a much larger scatter than any of the others, which is indicative that there may be RFI in this channel. Switching to the other calibrator source shows the same pattern. It is sometimes diagnostic to also examine the XY, YX polarizations here. For calibrator sources, we expect the polarized flux density to be a relatively small fraction of the unpolarized (i.e., XX, YY) one, whereas RFI is often strongly polarized. We conclude that channel ID=6 is likely affected by RFI, and therefore we will proceed to flag it also, once again using `flagdata`. Once again, we will flag that channel also for the target source data, since even if the target source data were not also affected by the RFI, the data would be uncalibratable.

CASA command:

```
flagdata(vis=msfile, spw='0:6', flagbackup=True)
```

We have now flagged most of the bad data affecting this run, in particular the bad data for the calibrator sources, and can now proceed to the actual calibration.

2.5 STEP 4... Calibration

In order to explain calibration, we must first explain briefly the nature of the visibility measurements we are trying to calibrate. If the field of view is small enough so that it can be represented as a plane (rather than as a curved part of the celestial sphere), then the sky brightness can be written as $I(l, m)$, where l and m are the two orthogonal coordinates (typically RA and decl). The visibility function measured by an ideal interferometer, V , is given by a two-dimensional Fourier transform of this sky brightness:

$$V(u, v) = \int \int I(l, m) e^{i2\pi(ul+vm)} dl dm \quad (2.1)$$

where u, v represent the coordinates in the Fourier transform plane. At any given instant, each baseline, or pair of telescopes, in the interferometer measures this visibility function at a point in the Fourier transform plane whose u, v coordinates are determined by the length and orientation of the baseline, as projected onto the sky plane. Since the telescopes rotate with the earth, each baseline typically sweeps out an elliptical path in the u, v plane as the earth rotates. Although I is

real, $V(u, v)$ is generally complex, with non-zero values for both real and imaginary parts. $V(u, v)$ is often described as an amplitude and a phase, rather than by its real and complex parts. Note that $V(u, v)$ is symmetric with $V(u, v) = V(-u, -v)$: generally the Fourier transform of a real function, like I , is complex but symmetric.

The correlator output differs from the V measured by an ideal interferometer for a variety of reasons, to do with both instrumental effects as well as propagation effects in the earth's atmosphere and ionosphere². The relationship between the observed and the true visibility on the baseline between two telescopes, i and j , can be written in a very general way as:

$$V_{ij}^{\text{observed}} = J_{ij} V_{ij}^{\text{true}} \quad (2.2)$$

J_{ij} represents the accumulation of all corruptions affecting the visibility measurement on baseline ij . Both V and J have complex values and are generally functions of frequency, polarization and time, but we have suppressed the additional indices pertaining to frequency, polarization and time for clarity.

Calibration involves determining the values of J_{ij} , in order to invert the above equation and to determine the best approximation possible of V_{ij}^{true} from the observed values of V_{ij}^{observed} .

This task of determining J_{ij} is made considerably easier by the fact that most of the effects contained in J_{ij} are *antenna*-based. In other words, J_{ij} can be factored into $J_i \otimes J_j^*$ (where J_j^* is the complex conjugate of J_j). For an array consisting of N_{ant} antennas there are therefore only N_{ant} unknown values of J_i which need to be determined, rather than the approximately N_{ant}^2 values of J_{ij} which would be required otherwise. Bear in mind that we have suppressed the time and frequency dependence here, so J_i is not a single complex number, but a function describing the complex response of antenna i in both time and frequency over the time-range and bandwidth of the observations.

The factorization occurs because the bulk of J_{ij} has its origin in effects which corrupt the radio signal as it travels through the atmosphere and the antenna electronics, these effects are therefore due to an individual antenna i , but do not depend on the other antenna j which makes up that particular baseline. Although there are baseline-dependent effects, they are mostly very small, and for most purposes, can be ignored.

It is practical to further factor the J_i into different components, representing different factors affecting the radio signal, each of which can be determined independently. Some of the components, for example, are time-invariant, and therefore need to be determined only once for the observing run, while others are time-variable. In our calibration process, we will first factor J_{ij} into the antenna-dependent J_i and

² At low frequencies, such as those observed by KAT-7, the effect of the ionosphere is usually the dominant one

thenceforth into a series of different factors, and each factor will be determined in a different step. Our approach is similar to the so-called “Measurement Equation”, which is a more formal approach to factoring J_{ij} described in [Sault and Cornwell, 1999], but we factor J_{ij} into slightly different components.

So we can write:

$$V_{ij}^{\text{observed}} = J_{ij} V_{ij}^{\text{true}} = (J_i \otimes J_j^*) V_{ij}^{\text{true}} \quad (2.3)$$

where

$$J_i = A_i B_i G_i \dots \quad (2.4)$$

For example, A_i might be corrections for some effect known a priori, such as an error in one of the antenna positions used in the correlator, and B_i might be to describe the bandpass response of the antennas, in other words the variation in sensitivity and phase between the different channels, which is taken to not vary with time, while G_i provides the variation in amplitude and phase response as a function of time, but taken not to be a function of frequency, i.e., determined for all frequency channels simultaneously. Note that the separation of these effects is not unique, however, we want to ensure that the final product J_i must be correct.

Some parts of J_i can be calculated without any reference to the visibility data. However, most of J_i will be determined from the values of V^{observed} for calibrator sources, where we know what the values of V^{true} should be and can therefore invert the equation above. In particular, most of the parts of J_i due to the atmosphere and ionosphere cannot be known in advance. The values of J_i (recall that J_i in general depend on time, but we have suppressed this dependence in the notation), determined from the observations of calibrator sources are then interpolated (or extrapolated) to the times corresponding to the observations of the intended target source(s). The various calibration programs perform the function of estimating the antenna-based calibration factors (e.g., A_i , B_i etc., from above) from the (baseline-based) calibrator source visibility measurements.

For example, the antenna-based bandpass response, B_i is typically treated as being independent of time (although it is by definition a function of observing frequency). It is typically determined by examining the visibilities for a strong calibrator source of known spectrum (the simplest case would be a source with a flat spectrum, with true flux density independent of the particular frequency or channel under consideration). Often this is the first step in the calibration, since once this bandpass correction is known, we can use an average across the channels to determine the remaining parts of J_i , rather than having to do it for each channel independently.

In order to use the visibility data from calibrator sources to calculate calibration quantities, such as the bandpass response, one should edit out the bad visibility data first. This is why we started above with examining the data and flagging bad points or antennas.

The typical sequence is to first use observations of a strong calibrator, often the flux-density calibrator, to determine the bandpass response. With that in hand, we no longer have to make separate solutions for every channel, but can combine them. One then uses the observations of calibrator sources to determine the variation of J_i with time. Many calibrator sources are variable on timescales of weeks or longer and their flux densities are therefore not accurately known. In order to set the flux density scale correctly, one needs observations of at least one flux density calibrator source, whose flux density *is* known. Typically, one uses one or a few scans of a flux density calibrator source for this purpose, and then the remainder of the J_i can be obtained from other calibrator sources whose flux density is not known (although it should not vary during the observing run), but which *are* known to be point-like. These secondary calibrators are known as phase calibrators or secondary amplitude calibrators.

So, let's begin the calibration. For safety and sanity, one should begin by “clearing the calibration”. The data structure of CASA measurement sets (ms) is that the observed data are stored in a DATA column of the ms, while the corrected data (i.e., those multiplied by J_{ij}) are stored in the CORRECTED_DATA column once they become available. The values from which J_{ij} are constructed are stored in various calibration tables. There is a further column of visibility data called MODEL_DATA where model values, if available, such as the expected visibility values for a calibrator of known structure, are stored. The task `clearcal` initializes the CORRECTED_DATA and MODEL_DATA columns and therefore clears calibrated data from any previous attempts at calibration from them.

CASA command:

```
clearcal(vis=msfile,field='')
```



Time to take a break and have some coffee... The steps after will need more attention.

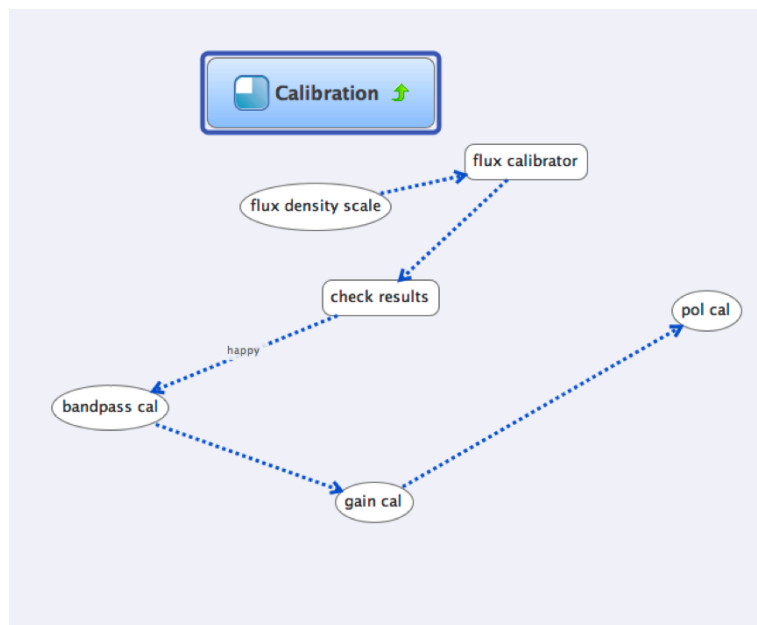


Figure 2.2: A schematic diagram of the calibration process.

There are few sources whose flux density is known reliably which can be used as primary amplitude calibrators. The task `setjy` knows about the most common of these, and will insert the correct flux density for the frequency of the observations as well as entering the model into the `MODEL_DATA` column for subsequent use.



These primary flux density calibrators are often somewhat resolved, and thus source models should be used to represent the source structure. The task `setjy` can be invoked to list the available source models (with `listmodels=True`). However, the flux density calibrator for our observations is PKS 1939-638, which is one that is close to ideal in that it has no structure on scales larger than $1''$. The default point model is therefore adequate and no further source model is required in our case.

We still have to enter the actual flux density for PKS 1939-638 for the actual frequency of the observations. In general the flux density varies with frequency, and `setjy` will calculate a slightly different value for each channel from the known spectrum. (In the case that a model image is used for the flux-density calibrator, `setjy` will also scale the model image to the total flux density for each channel in the ms, and then enter the Fourier transform of the scaled model image into the `MODEL_DATA` column for that source and channel.) Setting

fluxdensity = -1 causes setjy to look up the flux density for flux density calibrators and enter the standard value.

CASA command:

```
setjy(vis=msfile, field='PKS 1934-638',
      scalebychan=True, standard='Perley-Butler 2010',
      fluxdensity=-1)
```

setjy should produce the following output in the logger, from which one can see that it has identified calibrator source PKS 1934-638 and filled in the known flux density of ~ 13 Jy in I (the exact flux density will vary somewhat and be slightly different in the different channels, with the listed value being that for the channel at the reference frequency of 1932 MHz). If you see the message that your source “is not recognized by Perley-Butler 2010”, then setjy will have entered a default flux density of 1.0 Jy for your flux calibrator and you have to fix the problem otherwise your flux density scale could be wrong by a large factor.

CASA output:

```
setjy:::casa ##### Begin Task: setjy          #####
.....
setjy::imager::setjy() Using channel dependent flux densities
setjy::imager::data selection Selected 2709 out of 71904 rows.
setjy::imager::setjy() PKS 1934-638 (fld ind 0) spw 0 [I=13.182, Q=0, U=0, V=0] Jy, (Perley-Butler 2010)
....
```

Once we are happy that the correct flux density has been entered for the flux-density calibrator, we can proceed to the next step, which is to solve for the bandpass response, in other words for the variation of the complex antenna gains across the different channels. (Note that setjy above did not yet solve for any part of J_i , it only entered the flux density of our calibrator, essentially the value of V^{true} into the ms so that we *can* start solving for J_i in the next steps. The variation of the gains across the observing band is largely due to instrumental effects and usually quite stable, so we only need to solve for it once. Once we have solved for the bandpass response, we can combine all the channels in a single solution for the subsequent steps.

However, we have the problem that the phase part of J_i can be quite rapidly time-variable even though the amplitude part is not. In order to properly calculate the bandpass response, we therefore want to do some temporary phase-calibration for that scan which we intend to

use for the bandpass calibration, to prevent decorrelation when vector averaging the data in computing the bandpass solutions. In other words, to obtain the bandpass solution, we will average the complex visibility data in time over a whole scan. If there is any variation on phase over the scan, we want to calibrate that out before performing the bandpass solution.

We do this phase calibration using `gaincal`, and using only a portion of the band, in this case the six channels from 7 to 11. These solutions for the antenna phases as a function of time can then be used to calibrate out any variation in the phase during the scan for the purposes of the subsequent bandpass solution. The rapid phase variation we are calibrating out here are due mostly to the atmosphere, and are not expected to vary significantly between the channels. Therefore, although these phase solutions are determined from only a central portion of the band, they are applicable to the whole band (i.e., to all 19 channels). The choice of the particular channels (7 to 11) is relatively arbitrary, but usually its best to use a few channels near the centre of the band, which were found to be relatively free from interference in the editing steps earlier.

For this next step, we need to specify a reference antenna. When we ran `h5toms` we chose reference antenna 5. However, its possible that antenna 5 was not present at all during the run, or that we removed some of its data in flagging. You can use `plotms` to plot `antenna1` against `time` for all fields to quickly identify which antennas were present during which parts of the run. In this case you should see that antenna ID=0 is missing during the first part of the run, because we flagged it earlier. Antenna ID=5 seems to be present throughout, and can therefore serve as our reference antenna for the rest of the calibration. Note that it is possible to use different reference antennas for different parts of the calibration, but if possible you should use the same antenna throughout. It is also desirable to choose a reference antenna which is known to be stable and fairly sensitive (in the case of KAT-7 the sensitivities of the 7 antennas is very similar, so the only real criterion for picking the reference antenna is stability).

CASA command:

```
gaincal(vis=msfile, caltable=gain_table0,
        field='PKS 1934-638', refant=reference_antenna,
        spw='0:7~11', calmode='p', solint='int',
        minsnr=4, minblperant=4, solnorm=T, gaintable='')
```

Remember that we defined the variables `refant` and `gain_table0` earlier to hold the name of the reference antenna and gain table, respectively, that we're using for this particular ms.

After having determined the solutions, we should examine them using `plotcal`. The task `plotcal` can plot a variety of gain solutions in a variety of ways, check `help plotcal` to find out more. We remind you that `gaincal` has factored J_{ij} into antenna-based components, so you don't see a set of values for each baseline, but rather only one for each antenna (J_i), and the CASA programs will calculate the full value of J_{ij} as needed when calibrating the visibility values.

CASA command:

```
plotcal(caltable=gain_table0, xaxis='time',
        yaxis='phase', field='PKS 1934-638',
        iteration='antenna', plotrange=[0,0,-180,180])
```

On the `plotcal` plot, you should see that the gain-table phases (antenna phases, or the phase-component to J_i) should be fairly stable during the minute or two of the scan-length. Small variations are fine, indeed the point of running `gaincal` was exactly to find these small variations so that they can subsequently be calibrated out from the visibilities. However, if the antenna phase appears random from one point to the next then something has likely gone wrong.

Now we can run `bandpass` to solve for the complex gain as a function of channel across the observing band. Specifying `gaintable=[gain_table0]` causes `bandpass` to apply the time-dependent antenna phase calibration we determined in the previous step before forming the bandpass solutions.

CASA command:

```
bandpass(vis=msfile, caltable=bandpass_table0,
         field='PKS 1934-638', refant=reference_antenna,
         solnorm=True, combine='scan', solint='inf',
         bandtype='B', gaintable=[gain_table0])
```

You will see the message:

```
Insufficient unflagged antennas to proceed with this solve.
(time=2012/07/01/22:08:07.6 field=0 spw=0 chan=6)
```

Since we flagged channel 6 or `spw=0:6` earlier on, this is not something to worry about, its just telling us that there is no solution for that particular channel, but that doesn't matter because we have no visibility left for that channel. A message like this about some channel for which there *was* visibility data left would be a problem.

The task `plotcal` can also be used to display the bandpass solution:

CASA command:

```
plotcal(caltable=bandpass_table0, xaxis='chan',
        yaxis='amp', subplot=222, iteration='antenna')
```

First note that plots for the antennas for which the data have been flagged (`ant2`, `ant4`) are blank, which is as expected. For the others the gain amplitude varies only by 10% to 20% across the channels, which is in the expected range. Since we asked `bandpass` to normalize the solution, the average gain for each antenna is 1.0 — at this point we're only after the variation across the band. If any of the gains were to differ from unity by more than a factor of 2 or so, then may be more bad data to flag. One would then have to flag that bad data, and then re-start the calibrations from the `clearcal()` step.

The next step derive the main, time-variable, part of J_i , which we term G_i . This is done for each source individually, by comparing the observed visibilities V_{ij}^{observed} to the expected visibilities V_{ij}^{true} . For the flux density calibrator, we have determined the latter values, and we can get the correct values for G_i . For the phase calibrator sources, we assume a point source model, so we know the phase part of all the V_{true} is 0. We *assume* for now that the unknown flux-density of the phase-calibrator sources is exactly 1.0 Jy (typically they are of this order). If we then solve for G_i , we will obtain complex values whose amplitudes differ from the correct ones by a factor depending only on the true source flux density (by $([1 \text{ Jy}]/[\text{true source flux density}])^{0.5}$). We

can then compare the G_i obtained for the secondary calibrators, and scale them so that they match as well as possible the correctly-scaled amplitudes of G_i obtained for the primary flux-density calibrator. The scaling required gives us the true flux density of the secondary calibrator source in terms of the originally assumed one, i.e., 1 Jy. This process is known as “flux density bootstrapping”.

The reason for this complicated procedure is that there are only a handful of flux-density calibrators, so it is generally impossible to choose a flux density calibrator close to the target source on the sky. There are many phase calibrator sources, however, so a phase calibrator can be chosen that is much closer on the sky (typically a few degrees $^\circ$ away). The flux density bootstrapping allows us to use the phase and amplitude calibration (i.e. G_i) from these much nearer calibrator sources, which will provide a much better calibration for the target source, but to set the flux density scale accurately from observations of the flux density calibrator. The flux density scale is usually stable for hours, so once we have transferred the flux density scale to the phase calibrator sources by flux density bootstrapping, the solutions we obtain for the phase calibrator sources are equivalent to those which would have been obtained if we’d known the flux density of the phase-calibrators to begin with.

First we determine the values of G_i , also called the complex gains, for the flux-density calibrator:

CASA command:

```
gaincal(vis=msfile, caltable=gain_table1,
        field='PKS 1934-638', solint='inf',
        refant=reference_antenna, gaintype='G',
        calmode='ap', gaintable=[bandpass_table0])
```

Note that the a new solution is started for each scan, the `solint='inf'` only means that the solution interval can be arbitrarily long *up to* the next scan boundary.

Next we determine the complex gains for the phase calibrator (so far under the assumption that it has a flux density of 1 Jy).

CASA command:

```
gaincal(vis=msfile, caltable=gain_table1,
        field='PKS 1613-586', solint='inf',
        refant=reference_antenna, gaintype='G', calmode='ap',
        append=True, gaintable=[bandpass_table0])
```

Note that the `append=True` causes `gaincal` to append to the specified gain table, rather than over-writing it.

Plotting again to have a look at the solutions:

CASA command:

```
plotcal(caltable=gain_table1, xaxis='time',
        yaxis='amp', iteration='antenna')
```

Notice that, for each antenna, there are two sets of points. The lower set, in this case with gain amplitudes near ~ 0.1 corresponds to our flux density calibrator (PKS 1934-638). Since we were working with the correct flux density for this source (~ 13 Jy), these gains already have the correct values. The other set of points, with values near ~ 0.2 , are those for PKS 1613-586, which differ from the true values by a scale factor. The next step is to derive this scale factor, so as to make the two sets of points match as well as possible. The true flux density of PKS 1613-586 can be trivially determined from the derived scale factor.

The task `fluxscale` calculates the scale factor, thus performing the bootstrapping, and will output a correctly scaled version of the calibration table, which we can then use to calibrate the data, as well as printing out the derived value for PKS 1613-586's flux density.

CASA command:

```
fluxscale(vis=msfile, caltable=gain_table1,
          fluxtable=flux_table1, reference=['PKS 1934-638'],
          transfer=['PKS 1613-586'])
```


CASA output:

```

INFO fluxscale ##### Begin Task: fluxscale          #####
INFO fluxscale Opening MS: CirX1.ms for calibration.
INFO fluxscale Initializing nominal selection to the whole MS.
INFO fluxscale Beginning fluxscale--(MSSelection version)-----
INFO fluxscale Found reference field(s): PKS 1934-638
INFO fluxscale Found transfer field(s):  PKS 1613-586
INFO fluxscale Flux density for PKS 1613-586 in SpW=0 is: 4.6868 +/- 0.012777 (SNR = 366.816, N = 10)
INFO fluxscale Storing result in CirX1_spw0.fluxscale1

```

What is it we can deduce from this output? How do we know everything is fine? First, we can examine the messages produced by fluxscale. It has determined a flux density for PKS 1613-586 of 4.687 ± 0.013 Jy, making it a fairly strong source. For a strong source like this, the uncertainty in the bootstrapped flux density should be $< 1\%$, which indeed it is. You can also check that the flux density is within the range of values usually seen for your particular source at this frequency. For example, our source, PKS 1613-586, is listed in the ATCA calibrator list at (<http://www.narrabri.atnf.csiro.au/calibrators>) as having flux density of 4.21 Jy at 1.75 GHz (the closest of the listed frequencies to ours), so our value of 4.687 Jy is quite reasonable, as its only 16% higher. Variations of a factor of 2 are not uncommon, if you get a flux density more than a factor of 2 different than the listed, you may want to double check.

If you use plotcal as above, but on the new version of the table produced by fluxscale, i.e. caltable=flux_table1), you will see that the gain amplitudes for PKS 1913-638 are no longer readily distinguishable from those for PKS 1613-586, which was the goal (the gain of the antennas should not depend on which source you are observing!). You may also plot the gain phases with:

CASA command:

```

plotcal(caltable=flux_table1, xaxis='time',
        yaxis='phase', iteration='antenna',
        plorange=[0,0,-180,180])

```

And you should see that the gain phases are in fact very stable in time, and relatively consistent between our two calibrator sources. For ant5 the gain phases will be exactly 0° as this was our reference antenna, so it has a gain phase of 0° by definition. In order to calibrate the visibility data for our target source, we have to interpolate between the calibrator gain values. Since the gain phases for our phase

calibrator source PKS 1613-586 form a nice smooth line (in fact the individual points may not be distinguishable on your plot), the interpolation should yield an accurate value. If, on the other hand, the gain phases (or amplitudes) differ drastically from one point to the next, then interpolation to the intervening target-source observations will be dubious.

Our calibration therefore is looking good. The next step is to actually use it to calibrate the visibility data and produce our estimates of V_{ij}^{true} . Here is where we calculate the $V_{ij}^{\text{true}} = V_{ij}^{\text{observed}} J_{ij}^{-1}$. This is where we want to be careful about how we interpolate our gain solutions (J_i). For example, on the plot of gain phase against time, you could see the gain phases for PKS 1613-586 all lay on a fairly smooth curve, while those for PKS 1934-638 did not quite lie on the curve. Since PKS 1613-586 is much closer on the sky to our target source (Circinus X-1) the gain values (i.e., values of J_i) derived from PKS 1613-586 are better for interpolating to the target. However, we also want to calibrate the visibility data for PKS 1934-638, and of course here we're much better off using the gain solutions derived from PKS 1934-638 itself. (Note that although at this point, we're largely done with our calibrator sources, and will not have any immediate further use for the calibrator-source visibilities, it's probably a good idea to apply the calibration properly to the calibrator sources as well as the target source).

The easiest way to keep it straight is to apply the calibration separately to each source. The task that applies the calibration is called `applycal`. It combines the requested gain tables, interpolates or extrapolates as required to determine the full antenna-based J_i , calculates the baseline-based J_{ij} values as needed for each visibility measurement, and then writes out the properly calibrated visibility data (our estimate of V_{ij}^{true}) into the `CORRECTED_DATA` column of the ms. You can specify the `field`, which determines which source's visibility data get corrected, and `gainfield`, which specifies from which source's gain table entries the required correction is determined. The best calibration is usually when the `field` entry is the same as the `gainfield` entry, but usually one does not have calibration solutions for the target sources, so for the target source, `gainfield` must necessarily be different than `field`. Note that as `applycal` can simultaneously apply several calibration tables, `gainfield` can have multiple entries, one for each calibration table to be applied. The two calibration tables we want to apply are the bandpass table, `bandpass_table0`, and the gain table with the correct flux density scaling for PKS 1613-586 or `flux_table1`. Since there are two tables to apply, we need two entries in `gainfield`.

Lets just start with the first source: PKS 1934-638. We want to calibrate its visibility data and therefore set `field` to PKS 1934-638.

Since the `bandpass_table0` contains only entries for PKS 1934-638 anyway, there is no choice and we can just leave the corresponding `gainfield` entry blank – this will apply all sources in the table. For the `flux_table`, we want to use only the entries derived from PKS 1934-638 itself, so the second `gain_field` entry should be PKS 1934-638. The parameter `interp` specifies the nature of the interpolation used in gain table in a manner analogous to `gainfield` (see `help applycal` for details). In this case we will just choose the entry nearest in time to the particular visibility to be calibrated in `bandpass_table0` and interpolate linearly in `flux_table1`.

CASA command:

```
applycal(vis=msfile,
         gaintable=[bandpass_table0,flux_table1],
         field='PKS 1934-638',
         gainfield=['', 'PKS 1934-638'],
         interp=['nearest',''])
```

Next we turn to our phase calibrator. `field` should now be set to PKS 1613-586. Again we can leave the first entry in `gainfield` blank, since we again want to use any solutions in the bandpass table. However, this time we want the second `gainfield` entry to be PKS 1613-586.

CASA command:

```
applycal(vis=msfile,
         gaintable=[bandpass_table0,flux_table1],
         field='PKS 1613-586',
         gainfield=['', 'PKS 1613-586'],
         interp=['nearest',''])
```

Finally, we apply the calibration to the target field. Now `field` becomes Circinus X-1. The first entry of `gainfield` again can stay blank (we're again using the bandpass solution from PKS 1934-634). For the second entry in `gainfield`, however, we cannot use the same value as `field` like we did above, because we have no gain solutions for Circinus X-1. Since PKS 1613-586 is much closer on the sky, we want to use its entries from `flux_table1`, but not those from PKS 1934-634, so the second `gainfield` entry should again be PKS 1934-634.

CASA command:

```
applycal(vis=msfile,
         gaintable=[bandpass_table0,flux_table1],
         field='Circinus X-1',
         gainfield=['','PKS 1613-586'],
         interp=['nearest',''])
```

2.6 STEP 5... *Splitting the data*

We are now almost ready to split the calibrated visibility data for our target source into separate ms file. Note that this step is convenient, rather than necessary. One could just as well make images (and even proceed to self-calibration) using the original ms, but often it is convenient to split out the data of interest at this point. This particular data set has already been averaged in frequency to reduce the size for the purposes of the workshop. These days most raw data sets you encounter will likely have far more than 19 channels, and data sets from interferometers with more telescopes than KAT-7, such as JVL, ALMA, LOFAR and MeerKAT once it comes online, will have far more visibility measurements at each time-stamp. In such cases one might average down in frequency to reduce the size of the data set and make for faster processing. It is not necessary in this case, so our split output file will have the same number of channels as our original ms.

However, *before* you do any averaging, it is a good idea to briefly examine the calibrated visibility data for the target source in case there is more bad data we need to flag. Recall that so far we have only examined the data for our calibrator sources, but not yet that for Circinus X-1. Lets turn again to plotms. We will now plot correlated flux density against baseline length in wavelengths (uvwave). Such plots are commonly made to diagnose the nature of the source structure. We now have to specify that we want the CORRECTED_DATA column: we've gone to all this trouble to determine how to calibrate our visibilities, so we want to make sure we're enjoying the fruits of our labours.

CASA command:

```
plotms(vis=msfile, field='Circinus X-1',
       xaxis='uvwave', yaxis='amp',
       correlation='XX,YY', ydatacolumn='corrected',
       coloraxis = 'corr')
```

You should see something like this:




Figure 2.3: Plot of visibility amplitude against baseline-length for our target source, Circinus X-1.

This plot mostly shows that Circinus X-1 is *not* a point source, since the correlated flux density seems to vary systematically with baseline-length (uv -distance or uvw). There are, however, a few outlier points still, with anomalously high amplitudes above 10 Jy, such as the spikes at $UVwave \sim 140, 290 \lambda$. These are likely due to RFI. If we re-plot amplitude against channel, you will see that all the high amplitude point are in a single channel: channel 18. To further narrow it down, select `spw = '*:18'` and `xaxis='time'`. Not only do the high-amplitude points occur in one channel, they occur at a specific time: near 25:40. By zooming in on the plot using the “zoom” tool (select the magnifying glass icon, and use the mouse to draw a box on the screen showing the region you want to zoom in on. You can zoom back out using the “plot” button), to see that the bad data occurs between times of 25:32:30 and the end of the scan at 25:34:00. Although it is possible to flag graphically directly in `plotms`, I encourage you to again use `flagdata`. Firstly, it is better to positively identify the bad data to be flagged first, rather than zapping points at random on the screen. Secondly, using the `flagdata` command means that you can, if needed, precisely and accurately repeat the operation. Note also, that we could do more work to locate the bad data, it might for example be only one antenna or baseline which is bad. However, since its such a small amount of data (under 1.5 minutes in one channel only), its more expedient just to flag it all. Note also that `plotms` labels time axes in a slightly non-standard way with hours having values > 24 . For `flagdata` these values have to be translated into a more formal

timerange.

CASA command:

```
flagdata(vis=msfile,
         timerange='2012/07/02/01:32:30~2012/07/02/01:34:00',
         field='Circinus X-1', spw='0:18', flagbackup=True)
```

So now we've removed those outliers, we can split the data for our source of interest, Circinus X-1.

CASA command:

```
split(vis=msfile, outputvis='CirX1_split.ms',
      datacolumn='corrected', field='Circinus X-1')
```

2.7 STEP 6... CLEAN aka Deconvolution

For interferometric data, the (u,v) coverage is virtually never complete, and therefore the instrumental response, called the dirty beam or the point spread function, PSF, has some unwanted secondary responses, known as side lobes. The Fourier transform of the visibilities represents the convolution of the true image with the dirty beam.

There are various techniques to reduce the effects of the sidelobes, in other words to effect a deconvolution. In radio astronomy, deconvolution has been dominated by two non-linear algorithms: CLEAN and the Maximum Entropy Method (MEM).

The original formulation of CLEAN was by Högbom [1974], and worked purely in the image plane. We explain this method here as it is easier to understand. Most modern implementations of CLEAN actually work partly in the (u,v) plane to produce somewhat more accurate results than Högbom's original image-plane formulation, but are conceptually similar. CLEAN works iteratively, transferring flux density from the dirty image to a set of point-like "clean components", which represent the true distribution of emission on the sky or the deconvolved image. After each step, a residual image is formed which represents the original dirty image after subtraction of the CLEAN components so far. Initially, before the first CLEAN component is found and subtracted, the residual image is just the dirty image.

The Högbom CLEAN algorithm proceeds using these steps:

1. Find strength and position of the brightest point in the current residual image. CLEAN makes the plausible assumption that this

peak is mainly due to a real signal and only a minor part comes from the sidelobe responses to other real emission farther away in the image. Note that one should specify one or more “clean windows” in which to search for real emission.

2. Subtract a point source at this location. The brightness of the point source is taken to be the brightness of the brightest point multiplied by a factor $\gamma(\leq 1)$ called the loop gain. This point source is the CLEAN component, and it is added to the list of CLEAN components. The response of the instrument to a point source is the dirty beam pattern. Therefore, to subtract this point source from our dirty image, CLEAN subtracts a dirty beam pattern, scaled so that its peak is the same as that of the point source, and centred on the location of the brightest point, from the current residual image.
3. Go to step 1 unless any remaining point sources are below some level specified for convergence, or the maximum number of CLEAN components is reached.
4. Convolve the accumulated CLEAN components, which represent the model of the deconvolved emission, with an idealized “clean” beam (usually an elliptical Gaussian fitted to the central lobe of the dirty beam).
5. Add the final residual image to the convolution from the previous step. The result is what is called the CLEAN image. Note that for purposes of calibration, it is usually the CLEAN components that are used rather than the CLEAN image.

The CLEAN algorithm was designed for the case where the true brightness distribution contains only a few unresolved sources. It works well for cases where the true emission occurs only within small, well-separated, regions across the field of view, in other words if the field of view is *mostly* empty. The following excerpt has been taken from Robert Reid’s Thesis [?]. “Physics and filled aperture measurements give us several properties that images should have. Non-negativity: Astronomical objects, even black holes, cannot steal light away from our receivers, but raw interferometric images abound with negative patches. Requiring the image to be non-negative is a strong constraint on what should be used to fill in the gaps. Locality: Usually the range of directions from which light can be received is limited by the antennas being used. The rough location and extent of the source may also be known from observations at other wavelengths. Raw images always have ripples going off to infinity that should be quenched. If nowhere else, nulls in the antenna reception pattern are invalid places to receive light from. For many frequencies and parts of the sky, bright

radio sources are also quite rare and confined (as known from filled aperture measurements), so images are expected to be mostly empty. Smoothness: Sources are expected to be continuous, and usually to have continuous derivatives. Also, it is best to minimize the effect of the unmeasured part of the sampling plane on the final image. Perfectly sharp edges depend on the entire sampling plane, far beyond where interferometers can make measurements. Smoothing the image concentrates its region of dependence on the finite area that interferometers sample. Agreement with theory: At first this seems to throw out the baby with the bathwater, since true discoveries must be novel. True discoveries, however, must rest on true observations. Since deconvolution ideally agrees with all the measurements, we should not let non-measurements mislead us. Another way to see this is to think of the model as assisting us in separating reasonable interpolations from unreasonable ones for the image, but the parameters of the model itself are determined by the measurements. For example a circle has an infinite number of points on its circumference, but three points on its locus are enough to specify it, if we are sure that it is a circle. In any case, images tend to be constructed before models, thus this principle is not always applicable.”

Two main inputs to `clean`: the image size (`imsize`) and cell size (`cell`), which correspond to the size in pixels of the map, and the size in arcseconds of a pixel, respectively, must be carefully chosen. Here we show how to calculate these two parameters. By knowing the frequency of our observations we can calculate the primary beam (i.e., the area of the sky observed with the telescope).

First we choose the size of the pixels or cells. This choice is determined by the resolution of the interferometer, which is given by the synthesized beam. The value that is usually given is the full width at half-maximum (FWHM). Although the exact value for FWHM of the synthesized beam will depend on the details of the u, v -coverage and weighting for each particular observations, and approximate value for a given set of observations depends largely on the maximum baseline, B_{\max} , measured in wavelengths (λ) and is given by:

$$\text{FWHM}_{\text{synthesized beam}} \simeq \frac{1}{(B_{\max}/\lambda)} \text{radians, or} \quad (2.5)$$

$$\text{FWHM}_{\text{synthesized beam}} \simeq \frac{3438}{B_{\max}/\lambda} \text{arcminutes} \quad (2.6)$$

You can determine B_{\max} in λ by using `plotms` to plot the visibility amplitude as a function of the baseline length in λ or `uvwave`.

In order not to be limited by the pixel size, the central lobe of the dirty beam needs to be well resolved, and therefore a value of approximately $\frac{1}{4}$ of the above FWHM of the synthesized beam should

be chosen for the cell size. When CLEAN is run, one should check that the *smaller*, or minor axis of the elliptical fitted CLEAN beam is at least $3\times$ the size of the cells. If not, CLEAN should be stopped and restarted with better-suited values for `cell`.

In our case, we get an estimate of $3.7'$ for the resolution and thus $0.92'$ for the pixel size, which we will round down to $0.9'$ (note that since the requirement is to have pixels small enough to resolve the beam, rounding *down* is okay but not rounding up).

The other choice the user must make is the size of the region to be imaged. The maximum size of the region that it is sensible to image is determined by the field of view of the interferometer, which is just determined by that of individual antennas, and is independent of the array configuration. The FWHM of primary beam of a uniformly illuminated antenna is given by [Napier \[1999\]](#):

$$\text{FWHM}_{\text{primarybeam}} = \frac{1.02 \times \left(\frac{c}{\nu}\right)}{(\text{dishdiameter})} \text{radians} \quad (2.7)$$

In our case, the KAT-7 dishes have a diameter of 12 m, so at our observing frequency of 1822 MHz, the FWHM of the primary beam is $\sim 48'$. Of course, sources outside this field can still be seen by the interferometer, but at reduced amplitude. In particular, a source at a distance of FWHM/2 (remember that the FWHM is a diameter, not a radius) will be attenuated by 50% due to the primary beam response.

For the relatively coarse resolution offered by KAT-7, one can afford to image a relatively large region, and a reasonable choice is to image a region with $3\times$ the diameter of primary beam FWHM. In our case this is equal to $(3 \times 48' = 144')$. Since we have determined our pixel size to be $0.90'$, we need to image a region of $\frac{144'}{0.90'} = 160$. It is common to choose image sizes which are powers of 2, so we shall round this up to 256, which is the value we use for `imsize`.

When do I stop cleaning?

Stop cleaning when the residuals are noise like, and/or the clean has converged (the cleaned flux is no longer increasing)! Here are some of the parameters of `clean` which determine when it stops:

`niter` - Number of CLEAN iterations to do. This can be useful when you are doing tests, but this parameter has NO physical meaning. Instead set to large number and let either `threshold` or `do interactive` to stop the cleaning. Please also ensure you are not over cleaning your image. Have an eye on the `casapy log` that the total fluxes are not decreasing drastically.

`threshold` - Stop cleaning when peak residual has this value, give units (i.e. mJy). Ideally, one would like the image to have a noise approaching the theoretical one, in which case a `threshold` of $\sim 3\times$

the theoretical rms noise is appropriate. Note: to reach this limit the data must be well calibrated/flagged and suffer from no serious artifacts (resolved out extended structure/negative bowls, poor PSF/(u, v) coverage, limited dynamic range etc).

CLEANing deeply into the noise, that is with threshold set to a value below the image rms level is dangerous if you have a large CLEAN window. If the area of your CLEAN window is small, fewer than say 10 times the area of the CLEAN beam, then it is fairly safe to set threshold to some low value and just clean deeply. If the area of your CLEAN window is larger than overcleaning can be a problem. Basically, once you reach rms (whether close to theoretical or not), you are just picking noise up one place and putting it down in another. With a small CLEAN window, there is not much freedom here so nothing much can go wrong by moving a bit of noise around within your small window. With a large window strange artefacts can arise. Remember also, that rms in an otherwise blank region of your image represents only a *lower limit* to the uncertainty of the parts of your CLEAN image which have emission.

Starting the cleaning

Before you proceed to make an image, you will use the CASA viewer to look and perform various tasks. Viewer³ is a GUI for image manipulation.

³Follow the demo from: <http://casa.nrao.edu/CasaViewerDemo/casaViewerDemo.html>

CASA command:

```
clean(vis='CirX1_split.ms', imagename='CirX1_split.im',
      niter=5000, threshold='5.0mJy', psfmode='hogbom',
      interactive= True, imsize=[256,256],
      cell=['0.5arcmin' , '0.5arcmin'], stokes='I',
      weighting='briggs', robust=0.0)
```

3

Spectral line calibration and imaging

Spectral line calibration is pretty similar to that for continuum. The real differences come in when imaging the data. Instead of averaging together all the data before imaging, we make one image per frequency channel, producing a **data cube**. In this tutorial we assume that you have worked through Chapter 3 and understand the basics of CASA and calibration.

The KAT-7 spectral line modes produce 4096 channels of data. The outer 500 channels are usually discarded since the gain tapers off rapidly as part of the guard bands. For this tutorial we will use a subsection of data to speed things up. (The original data file is 34 GB).

We will be looking at a massive star forming cloud in the Galactic Plane. At L-band we see a number of evolved HII regions (continuum) and hydroxyl masers (spectral lines). The msfile is called G330_OH.ms.

3.1 Inspection and flagging

First, let us set up all our filenames:

CASA command:

```
prefix = 'G330_OH'
msfile = prefix+'.ms'
gtable0 = prefix + '.G0'
btable0 = prefix + '.B0'
gtable1 = prefix + '.G1'
ftable1 = prefix + '.fluxscale1'
```

Lets have a look at what we have here:

CASA command:

```
listobs(msfile)
```

CASA output:

```
#####
#### Begin Task: listobs      ####
listobs(vis="G330_OH.ms",selectdata=True,spw="",field="",
        antenna="",uvrange="",timerange="",correlation="",scan="",
        intent="",feed="",array="",observation="",verbose=True,
        listfile="",listunfl=False,cachesize=50)
=====
MeasurementSet Name: /home/sharmila/spec_line_tut_OH/G330_OH.ms      MS Version 2
=====
Observer: sharmila      Project: 20130621-0006
Observation: KAT-7
Data records: 12327      Total integration time = 40778.8 seconds
Observed from 21-Jun-2013/14:08:01.0 to 22-Jun-2013/01:27:39.8 (UTC)

ObservationID = 0      ArrayID = 0
Date      Timerange (UTC)      Scan      FldId      FieldName      nRows      SpwIds      AveInts
21-Jun-2013/14:07:45.9 - 14:12:17.7      1      0      1934-638      189      [0]      [30.2]
      14:25:02.7 - 14:25:32.9      4      1      1613-586      21      [0]      [30.2]
      14:26:13.2 - 14:30:45.0      5      2      G330.89-0.36      189      [0]      [30.2]
      14:36:37.3 - 14:37:07.5      7      1      1613-586      21      [0]      [30.2]
      14:37:47.8 - 14:42:19.6      8      2      G330.89-0.36      189      [0]      [30.2]
      14:48:11.9 - 14:48:42.1      10      1      1613-586      21      [0]      [30.2]
      14:49:22.3 - 14:53:54.1      11      2      G330.89-0.36      189      [0]      [30.2]
      .....
      .....
      .....
      01:10:33.0 - 01:11:33.4      159      1      1613-586      42      [0]      [30.2]
      01:11:48.5 - 01:16:20.3      160      2      G330.89-0.36      189      [0]      [30.2]
      01:22:07.6 - 01:22:37.8      162      1      1613-586      21      [0]      [30.2]
      01:23:23.1 - 01:27:54.9      163      2      G330.89-0.36      189      [0]      [30.2]
      (nRows = Total number of rows per scan)

Fields: 3
ID      Code      Name      RA      Decl      Epoch      nRows
0      T      1934-638      19:39:25.017468 -63.42.45.60158 J2000      2079
1      T      1613-586      16:17:17.893107 -58.48.07.88902 J2000      1176
2      T      G330.89-0.36      16:10:20.541228 -52.06.14.90063 J2000      9072

Spectral Windows: (1 unique spectral windows and 1 unique polarization setups)
SpwID      Name      #Chans      Frame      Ch1(MHz)      ChanWid(kHz)      TotBW(kHz)      Corrs
0      none      1001      TOP0      1665.985      0.381      381.9      XX YY

The SOURCE table is empty: see the FIELD table
Antennas: 7:
ID      Name      Station      Diam.      Long.      Lat.
0      ant1      ant1      12.0 m      +021.24.39.4 -30.33.10.2
1      ant2      ant2      12.0 m      +021.24.41.9 -30.33.09.1
2      ant3      ant3      12.0 m      +021.24.38.6 -30.33.09.1
3      ant4      ant4      12.0 m      +021.24.37.7 -30.33.09.1
4      ant5      ant5      12.0 m      +021.24.37.1 -30.33.10.0
5      ant6      ant6      12.0 m      +021.24.36.2 -30.33.12.5
6      ant7      ant7      12.0 m      +021.24.35.2 -30.33.07.5
#### End Task: listobs      ####
#####
```

We have averaged to 30 second integrations and there are 601 channels in this dataset to reduce the size of the file.

Now we set up our calibrator variables. We use PKS 1934-638 as the flux and bandpass calibrator and PKS 1613-586 as the gain calibrator. Antenna 6 is usually very stable, so we will use it as our reference antenna.

CASA command:

```
f_cal = '1934-638'
b_cal = '1934-638'
g_cal = '1613-586'
ref_ant = 'ant6'
```

Before we start calibration, it is vital to make sure that our calibrators are free of RFI. First we check along the frequency axis. This is a very narrow-band observation (1.1 MHz) so we don't usually expect to see much RFI. However, we have found self-generated RFI on some antennas, so we will look at each baseline independently.

CASA command:

```
plotms(vis = msfile,
       field = f_cal,
       xaxis = 'channel',
       yaxis = 'amp',
       iteraxis = 'baseline',
       yselfscale = True)
```

Did you spot the RFI? Identify the channel and use `flagdata` to flag the channel for the affected antennas.

CASA command:

```
flagdata(vis = msfile,
        spw = '0:548',
        antenna = 'ant5,ant6,ant7')
```

3.2 Set up models for the calibrators

Now let's prepare the file for calibration. This step will clear all previous calibrations or, if it is a freshly made ms file, will add in the

MODEL and CORRECTED DATA columns.

CASA command:

```
clearcal(msfile)
```

Now we are going to set up models for our calibrators. First, we set up our flux calibrator:

CASA command:

```
setjy(vis = msfile,
      field = f_cal,
      fluxdensity = -1,
      standard = 'Perley-Taylor 99')
```

PKS 1613-586 is not a very good calibrator, but it is difficult to find isolated sources close to the Galactic Plane. Instead, we use a model of the field, found by imaging the calibrator. You should have the directory 1613-586.model in your working directory. This is a CASA model file, which is produced by CLEAN. This will populate the MODEL column for the source with the appropriate visibilities.

CASA command:

```
setjy(vis = msfile,
      field = '1613-586',
      fluxdensity = [0,0,0,0],
      modimage = '1613-586.model')
```

3.3 Calibration

We first do a preliminary time-dependent phase calibration over a subset of the band. We have to find a small section of the band that is reasonably flat, and have enough bandwidth to achieve a reasonable signal-to-noise. The number of channels is going to depend on the total bandwidth of the observation. In this case, we are using $1/4$ of the total channels in the 1.5 MHz mode and don't expect to see much change across the bandpass.

CASA command:

```
plotms(vis=msfile,
       field = f_cal,
       xaxis = 'channel',
       yaxis = 'phase',
       avgtime = '1e8',
       correlation = 'XX',
       antenna = 'ant6',
       coloraxis = 'antenna2')
```

You will have noticed in `plotms` that the bandpass appears to be relatively flat across this channel range, so we can in fact use all of the channels. Note that this is not the case for the wider HI modes. You may find solutions failing if you use too large a channel range.

CASA command:

```
ref_chans = '0:1~600'
gaincal(vis = msfile,
       field = b_cal,
       caltable = gtable0,
       refant=ref_ant,
       spw = ref_chans,
       calmode = 'p',
       solint = 'inf',
       minsnr = 5,
       solnorm = True,
       interp = 'nearest')
```

Inspect the solutions using `plotcal`.

CASA command:

```
plotcal(caltable = gtable0,
       xaxis = 'time',
       yaxis = 'phase',
       markersize=3,
       plotsymbol='.',
       iteration='antenna',
       subplot=421,
       fontsize=8)
```

All looks good, so let's continue on to the bandpass calibration. We have found that in general, we do not need to do a time-dependent bandpass calibration for KAT-7, even though we visit the bandpass calibrator every hour in these observations (just in case), so we average over all of the bandpass scans to increase our signal-to-noise ratio. To start with, we are going to use the normal 'per channel' bandpass solution ie `bandtype='B'`. However, if the solutions turn out to be very noisy, it may be better to use a polynomial fit to the bandpass `bandtype='BPOLY'`

CASA command:

```
bandpass(vis = msfile,
         caltable = btable0,
         field = b_cal,
         refant = ref_ant,
         solnorm = True,
         combine = 'scan',
         solint = 'inf',
         bandtype = 'B',
         minsnr = 5,
         gaintable = [gtable0],
         interp = ['nearest'])
```

If you are using CASA 4.1 and above, there is a new task to plot bandpasses. This will enable you to see both the amplitude and phase, and overplot polynomial solutions (Figure 3.1).

CASA command:

```
plotbandpass(caltable = btable0,
             xaxis = 'freq',
             yaxis = 'both',
             subplot = 42)
```

The problem with such narrow band observations is that we have increased noise. Note that the peak-to-peak amplitude variation is about 10% and the phase variation is about 8 degrees. However, there is not much variation in the bandpass, so we can use a polynomial-based bandpass calibration. Note below that we change `bandtype` to `'BPOLY'`. We can specify different orders of polynomial fit for the phase and amplitude if necessary (Figure 3.2).

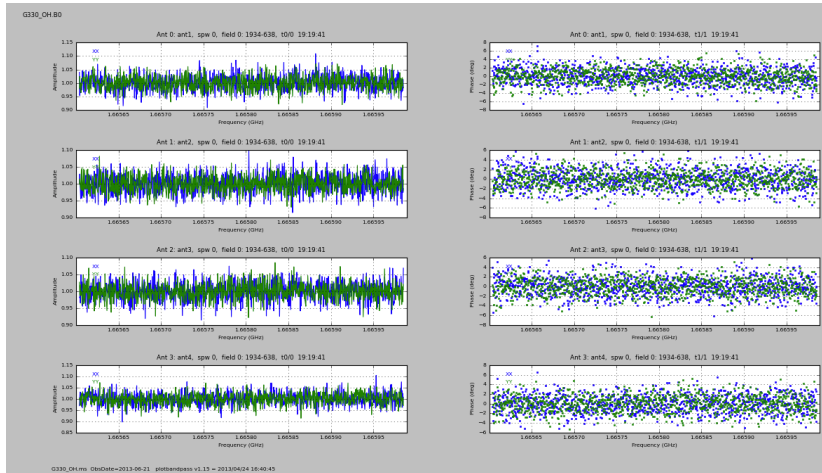


Figure 3.1: The output of plotbandpass for a few antennas.

CASA command:

```
btable1 = prefix + '.B1'
bandpass(vis=msfile,
         caltable = btable1,
         field = b_cal,
         refant = ref_ant,
         solnorm = True,
         combine='scan',
         solint = 'inf',
         bandtype = 'BPOLY',
         degamp = 10,
         degphase = 10,
         minsnr = 5,
         gaintable = [gtable0],
         interp = ['nearest'])
```

Why is the bandpass calibration so important? In spectral line observations we are interested in source structure as a function of frequency, so we want to remove any spurious variations due to the instrumental response. We may also want to subtract off continuum emission in order to simplify our imaging and in this case we don't want there to be any residual continuum emission. How do we know if we have sufficient signal to noise in our bandpass calibrator? We typically want the rms noise from the calibrator to be much less than the noise in the source. We use calibrators with reasonably high flux density, but spend less time on them than we do on our target sources,




Figure 3.2: The output of plotbandpass with B and BPOLY solutions overlaid.

so its good to review whether your observations satisfy the criterion that $I_{cal}\sqrt{t_{cal}}$ should be significantly greater than $I_{target}\sqrt{t_{target}}$. In the case of this observation, the target actually has quite a high continuum flux, and it will not be possible to meet this requirement without spending an impractical amount of time on the bandpass calibrator.

Now we calibrate the gain on our flux calibrator, which in this observation happens to also be our bandpass calibrator.

CASA command:

```
gaincal(vis = msfile,
        caltable = gtable1,
        field = f_cal,
        solint = 'inf',
        refant = ref_ant,
        gaintype = 'G',
        calmode = 'ap',
        solnorm = False,
        minsnr = 5,
        gaintable = [btable1],
        interp = ['nearest'])
```

Now check that there are no bad outliers in the gain calibration.

CASA command:

```

plotcal(caltable = gtable1,
        field = f_cal,
        xaxis = 'time',
        yaxis = 'amp',
        markersize=3,
        plotsymbol='.',
        iteration='antenna',
        subplot=421,
        fontsize=8)

plotcal(caltable = gtable1,
        field = f_cal,
        xaxis = 'time',
        yaxis = 'phase',
        markersize=3,
        plotsymbol='.',
        iteration='antenna',
        subplot=421,
        fontsize=8)

```

Calibrate the gains on the phase calibrator. Note that we append the solutions to the table containing the gains on the flux calibrator.

CASA command:

```

gaincal(vis = msfile,
        caltable = gtable1,
        field = g_cal,
        solint = 'inf',
        refant = ref_ant,
        gaintype = 'G',
        calmode = 'ap',
        append = True,
        solnorm = False,
        minsnr = 5,
        gaintable = [btable1],
        interp = ['nearest'])

```

As usual, check the gain solutions and make sure that there's nothing nasty going on. Next, we transfer the flux scale from our flux calibrator to the gain calibrator. This will automatically transfer the

flux calibration to the target when we apply the calibrations.

CASA command:

```
fluxscale(vis=msfile,
          caltable = gtable1,
          fluxtable = ftable1,
          reference = f_cal,
          transfer = g_cal)
```

CASA output:

```
#####
##### Begin Task: fluxscale #####
fluxscale(vis="G330_OH.ms",caltable="G330_OH.G1",
          fluxtable="G330_OH.fluxscale1",reference="1934-638",
          transfer="1613-586",listfile="",append=False,refspwmap=[-1],
          incremental=False, fitorder=1)
Opening MS: G330_OH.ms for calibration.
Initializing nominal selection to the whole MS.
Beginning fluxscale--(MSSelection version)-----
Found reference field(s): 1934-638
Found transfer field(s): 1613-586
Flux density for 1613-586 in SpW=0 is: 1.01027 +/- 0.00741467 (SNR = 136.252, N = 14)
Storing result in G330_OH.fluxscale1
Writing solutions to table: G330_OH.fluxscale1
##### End Task: fluxscale #####
#####
```

Check that the derived flux is close to what is expected. Now we are ready to apply the calibration and get ready for imaging.

CASA command:

```
target = 'G330.89-0.36'
applycal(vis=msfile,
         field = target,
         gaintable = [btable1, ftable1],
         gainfield = [b_cal, g_cal],
         interp = ['nearest', 'linear'])
```

CASA command:

```
split_outputvis = target+'.split.ms'
split(vis = msfile,
      outputvis = split_outputvis,
      field = target,
      datacolumn = 'corrected')
```

3.4 *Velocity rest frames*

There are a few things that you need to do before you can dive in and start cleaning your image. Since spectral line observations give us a lot of information about velocities of the objects that we are observing, it is important to be clear on what reference frames we are using. The apparent frequency of a source, also known as the 'sky frequency', is influenced by Earth's rotation and revolution around the Sun, the motion of our Solar System around the Galaxy, the movement of our Galaxy within the Local Group and there is also motion against the cosmic microwave background. The various rest frames are listed in Table 3.1. At L-band, the sky frequency change by up to 0.15 MHz during the course of a year.

The KAT-7 system does not do any Doppler corrections, so its rest frame is **topocentric**. The actual sky frequency of your source will vary from day to day and if you are observing with high frequency resolution, you may even see your spectral line shifting during the course of a single day's observation. This is definitely the case for these hydroxyl maser observations. We have a velocity resolution of 68 m/s so we do see the source Doppler-shifted by Earth's rotation.

You will usually want to have your final data cubes output into a different rest frame. The velocity frame to use depends on your science. In the case of this dataset, we are looking at a Galactic star formation region. It is meaningful in this case to work in the **Local Standard of Rest (LSR) frame**. There are actually two definitions of LSR, the kinematic LSR (LSRK) and dynamic LSR (LSRD). Most of the time LSRK is used and is generally synonymous with LSR. The other commonly used rest frame is **barycentric**, and has mostly replaced the old heliocentric standard.

There are two ways to change rest frames in CASA. The CLEAN task can change the reference frame of the spectral axis which is fine if you don't have to worry about doppler shifts during the course of the observation. CVEL is a more general spectral regridding tool, which enables you to correct for Doppler shifts, or change the channelisation

on your data prior to imaging. Doing the spectral regridding before running CLEAN can save you time if you think you are going to be trying different imaging settings.

Rest Frame Name	Rest frame	Corrects for	Max (km/s)
Topocentric	Telescope	Nothing	0
Geocentric	Earth Centre	Earth Rotation	0.5
Earth-Moon Barycentric	Earth-Moon centre of mass	Motion about Earth + Moon centre of mass	0.013
Heliocentric	Centre of Sun	Earth orbital motion	30
Barycentric	Earth+Sun centre of mass	Earth+Sun centre of mass	0.012
Local standard of rest (LSR)	Centre of mass of local stars	Solar motion rel to nearby stars	20
Galactocentric	Centre of Milky Way	Milky Way rotation	230
Local Group Barycentric	Local Group centre of mass	Milky Way motion	100
Virgocentric	Centre of the Local Virgo Supercluster	Local Group motion	300
Cosmic Microwave Background	CMB	Local Supercluster motion	600

Table 3.1: Velocity rest frames



You can use the viewer to look at visibilities. Lets try it now.

Figure 3.3 shows a pretty severe shift over the course of the observations. If we had to image this dataset without correcting for the Doppler shift, we would end up blending our spectral features together. We are going to be using CVEL to correct for the doppler shift, and at the same time, change the velocity reference frame to the Local Standard of Rest (V_{LSR}). Before we do that, we need to set the line rest

frequency in the measurement set. The KAT-7 observation framework does not make provision for setting the rest frequency in the data files so we shall have to do it manually. The rest frequency for this transition of OH is 1665.40184 MHz. We are going to open the spectral window table and insert the rest frequency. The default rest frequency in the file is usually the central observing frequency. Remember to close the file before continuing because it will be locked until you do.




Figure 3.3: A raster plot of visibilities on a single baseline. Note how the spectral line is shifting through the channels over time.

CASA command:

```
restfreq=1665.40184e6
tb.open(split_outputvis+'/SPECTRAL_WINDOW',
        nomodify=False)
tb.putcell('REF_FREQUENCY', 0, restfreq)
tb.close()
```

If run `listobs` now you should see that the reference frequency has changed. Next we run `cvel`.

CASA command:

```
freq_string = str(restfreq/1e6)+'MHz'
cvel_outputvis = target+'.cvel.ms'
cvel(vis = split_outputvis,
      outputvis = cvel_outputvis,
      mode = 'velocity',
      interpolation = 'linear',
      outframe = 'LSRK',
      restfreq = freq_string)
```




Figure 3.4: A raster plot of visibilities on a single baseline after doppler correction.

This can take some time for a large dataset. If you find yourself running low on disk space, you can now delete the split file since we will use the cvel file from now on. Have a look at the spectral lines using the viewer and convince yourself that we have removed the doppler shift. In Figure 3.4 you will see that the spectral lines now stay in the same channel. Note that some of the edge channels are now missing data. Make sure that you have split off sufficient channels to account for this before you apply cvel. We will lose about 10 channels in this process. When you clean your datacube, you should make sure that these channels are not included. I generally avoid this problem by

setting the start velocity and number of channels for the output from `cvel` such that I cover only the channels that I want to image. This trick works well if you have observations spread over more than a month, because the channel in which a particular velocity will be present will vary over the year. This brings your data into a common reference frame with consistent channel numbers. If you don't know what these parameters should be, you can run `cvel` on the full channel range, then examine the resulting file and rerun `cvel` with a smaller output range once you know which velocity range is of interest. Remember to leave sufficient line-free channels on either side of your dataset for continuum subtraction.

3.5 Continuum subtraction

This particular field is of a massive star formation region. The one degree field of view includes several evolved HII regions as well as younger regions containing masers. It will simplify the cleaning process considerably if we can subtract the continuum and image it separately from the line emission. Lets have a look at a spectrum and see if we can establish the continuum baseline. In the case of these masers, the emission is so strong that we can see it in individual integrations, however you need to be careful because there are also weaker features less than one Jy. In the case of weaker sources, we would need to average together all of the data in order to spot the emission. This works well only if the source is close to the phase centre and not too extended. The best way is to make a dirty cube and examine this through the viewer in order to determine line channels. In the case of extremely weak line sources you may even need to do some cleaning first before you can identify the lines.

Lets have a quick look at the averaged spectrum.

CASA command:

```
plotms(vis=cvel_outputvis,
       ydatacolumn = 'data',
       xaxis = 'channel',
       yaxis = 'amp',
       avgtime='1e8',
       avgscan=True)
```

There may be some emission around the last channels. It also turns out that this plot does not show sufficient detail to figure out where the line emission is. It will be much better to make a dirty cube first

and inspect the result to determine where we can place the continuum windows. To make a dirty cube, we simply set the number of iterations of CLEAN to zero.

CASA command:

```
clean(vis = cvel_outputvis,
      imagename = 'dirty',
      mode = 'channel',
      outframe = 'LSRK',
      restfreq = freq_string,
      imsize = 256,
      cell='30arcsec',
      stokes = 'I',
      niter = 0,
      psfmode = 'hogbom',
      interactive = False,
      weighting = 'briggs')

viewer('dirty.image')
```

We have a reduced channel range in this dataset in order to keep the file sizes down, and as a result most of the channels contain line emission. Continuum subtraction is a lot easier in the original 4096 channel dataset. Careful examination of the cube will show that the only line-free channels are around 280 to 340. The continuum subtraction task will issue a warning about extrapolating over the outer frequency ranges but will still do it.

CASA command:

```
cont_window = '0:280~340'
uvcontsub(vis=cvel_outputvis,
          fitspw = cont_window,
          fitorder = 1,
          want_cont = True)
```

Now you have two new measurement sets, one containing the continuum data, which you can image as usual, and one containing the spectral data. We do not expect to see any absorption features in this field, so if you do see negative features, it may mean that the continuum subtraction has failed. You can check this before going on to the laborious task of deconvolution by making a dirty cube of the

continuum-subtracted data.

3.6 *Finally, on to imaging*

Now we can start cleaning the spectral line cube. Masers typically have angular sizes of the order of milli-arcseconds and therefore are not resolved with KAT-7. However, there are five individual sources of maser emission in this field of view. Lets start the cleaning process. You will notice that we still specify the rest frequency even though we have already set it in the measurement set. This is because of some inconsistencies in the way CASA handles spectral velocities - for now it is safer to always specify the rest frequencies and the required output frame.

CASA command:

```
restfreq=1665.40184e6
freq_string = str(restfreq/1e6)+'MHz'
contsub_vis = target+'.cvel.ms.contsub'
cube_namebase = target+'.cube.clean'

clean(vis = contsub_vis,
      imagename = cube_namebase,
      spw = '0:1~590',
      mode = 'channel',
      outframe = 'LSRK',
      restfreq = freq_string,
      imsize = 256,
      cell='30arcsec',
      stokes = 'I',
      threshold = '0.15Jy',
      niter = 20000,
      psfmode = 'hogbom',
      interactive = True,
      weighting = 'briggs')
```

If you find you are running out of time, try setting interactive to False and let CASA find its own way.

Optional: Try cleaning the continuum. Below are the commands to start the continuum cleaning. Note that I have dropped the last 11 channels to get rid of the empty channels that we were left with after CVEL.

CASA command:

```

cont_vis = target+'.cvel.ms.cont'
cont_image = target+'.cont.clean'

clean(vis=cont_vis,
      imagename = cont_image,
      spw = '0:0~590',
      imsize = 512,
      cell='30arcsec',
      stokes = 'I',
      threshold = '10mJy',
      niter = 5000,
      psfmode = 'hogbom',
      interactive = True,
      weighting = 'briggs')

```

You will notice that this field has a pretty complex structure, consisting of point sources as well as more nebulous evolved HII regions. It is far better to clean this image using multi-scale clean, which will fit for several size-scales. To set up multi-scale, we need to specify the scales, in number of pixels, where 0 indicates a point source, 5 pixels is about one beam, and then we do a few multiples of the beam, going up to 4 times the beam in this case. Note that CASA will complain if the scales are too large and then drop them. If doing the cleaning non-interactively, it is good to also tell CASA to stop when it starts creating too many negative components, otherwise things can go wrong very quickly.



Warning: multiscale clean is a lot slower than the conventional clean because it repeats the clean cycle for every scale specified and needs to re-scale the psf for every scale. When setting this task to non-interactive mode, and trying to clean down to 5 times the expected rms noise, this takes almost 3 hours on a high-end desktop machine. This is also why we try to subtract the continuum off in the uv-plane, since it speeds up imaging the OH masers considerably.

CASA command:

```
cont_image = target+'.cont.multiscale.clean'  
clean(vis=cont_vis,  
      imagename = cont_image,  
      multiscale = [0, 5, 10, 15, 20],  
      negcomponent = 1,  
      spw = '0:0~590',  
      imsize = 512,  
      cell='30arcsec',  
      stokes = 'I',  
      threshold = '10mJy',  
      niter = 5000,  
      psfmode = 'hogbom',  
      interactive = True,  
      weighting = 'briggs')
```


4

Troubleshooting and Support

4.1 Getting help

In casapy, you can type `help [commandname]` to get help in a particular command.

If you've encountered a problem, have a question, or would like to report a bug, please send an email to our mailing list or visit our website.

Bibliography

J. A. Högbom. Aperture Synthesis with a Non-Regular Distribution of Interferometer Baselines. *Astron. Astrophys. Suppl. Ser.*, 15:417, June 1974.

P. J. Napier. The Primary Antenna Elements. In G. B. Taylor, C. L. Carilli, and R. A. Perley, editors, *Synthesis Imaging in Radio Astronomy II*, volume 180 of *Astronomical Society of the Pacific Conference Series*, page 37, 1999.

R. J. Sault and T. J. Cornwell. The Hamaker-Bregman-Sault Measurement Equation. In G. B. Taylor, C. L. Carilli, and R. A. Perley, editors, *ASP Conf. Ser. 180: Synthesis Imaging in Radio Astronomy II*, pages 657–669, 1999.